

Introduction to Quantitative Text Analysis (QTA)

Day Two Lab

Brian Fogarty

15 May 2018

Whereas last week's lab materials were about getting our text data set-up, etc., this week is about different statistical analysis options that are available with *Quanteda*.

1. Set-Up

This is about just getting set-up. More details at <https://tutorials.quanteda.io/introduction/install/>. In today's lab, we are going to work with the full inaugural speeches corpus and the 2010 Irish budget speeches data that is part of the *Quanteda* package.

```
setwd("C:/QTA/Day Two")
getwd()
```

```
## [1] "C:/QTA/Day Two"
```

```
.libPaths("C:/Program Files/R/R-3.4.2/library")
.libPaths()
```

```
## [1] "C:/Program Files/R/R-3.4.2/library"
```

```
library(quanteda)
```

```
## Warning: package 'quanteda' was built under R version 3.4.4
```

```
## Package version: 1.2.0
```

```
## Parallel computing: 2 of 4 threads used.
```

```
## See https://quanteda.io for tutorials and examples.
```

```
##
```

```
## Attaching package: 'quanteda'
```

```
## The following object is masked from 'package:utils':
```

```
##
```

```
##      View
```

```
library(readtext)
```

```
full_inaug_corp <- corpus(data_corpus_inaugural)
ndoc(full_inaug_corp)
```

```
## [1] 58
```

```
summary(full_inaug_corp)
```

```
## Corpus consisting of 58 documents:
```

```
##
```

	Text	Types	Tokens	Sentences	Year	President	FirstName
##	1789-Washington	625	1538	23	1789	Washington	George
##	1793-Washington	96	147	4	1793	Washington	George

##	1797-Adams	826	2578	37	1797	Adams	John
##	1801-Jefferson	717	1927	41	1801	Jefferson	Thomas
##	1805-Jefferson	804	2381	45	1805	Jefferson	Thomas
##	1809-Madison	535	1263	21	1809	Madison	James
##	1813-Madison	541	1302	33	1813	Madison	James
##	1817-Monroe	1040	3680	121	1817	Monroe	James
##	1821-Monroe	1259	4886	129	1821	Monroe	James
##	1825-Adams	1003	3152	74	1825	Adams	John Quincy
##	1829-Jackson	517	1210	25	1829	Jackson	Andrew
##	1833-Jackson	499	1269	29	1833	Jackson	Andrew
##	1837-VanBuren	1315	4165	95	1837	Van Buren	Martin
##	1841-Harrison	1896	9144	210	1841	Harrison	William Henry
##	1845-Polk	1334	5193	153	1845	Polk	James Knox
##	1849-Taylor	496	1179	22	1849	Taylor	Zachary
##	1853-Pierce	1165	3641	104	1853	Pierce	Franklin
##	1857-Buchanan	945	3086	89	1857	Buchanan	James
##	1861-Lincoln	1075	4006	135	1861	Lincoln	Abraham
##	1865-Lincoln	360	776	26	1865	Lincoln	Abraham
##	1869-Grant	485	1235	40	1869	Grant	Ulysses S.
##	1873-Grant	552	1475	43	1873	Grant	Ulysses S.
##	1877-Hayes	831	2716	59	1877	Hayes	Rutherford B.
##	1881-Garfield	1021	3212	111	1881	Garfield	James A.
##	1885-Cleveland	676	1820	44	1885	Cleveland	Grover
##	1889-Harrison	1352	4722	157	1889	Harrison	Benjamin
##	1893-Cleveland	821	2125	58	1893	Cleveland	Grover
##	1897-McKinley	1232	4361	130	1897	McKinley	William
##	1901-McKinley	854	2437	100	1901	McKinley	William
##	1905-Roosevelt	404	1079	33	1905	Roosevelt	Theodore
##	1909-Taft	1437	5822	159	1909	Taft	William Howard
##	1913-Wilson	658	1882	68	1913	Wilson	Woodrow
##	1917-Wilson	549	1656	59	1917	Wilson	Woodrow
##	1921-Harding	1169	3721	148	1921	Harding	Warren G.
##	1925-Coolidge	1220	4440	196	1925	Coolidge	Calvin
##	1929-Hoover	1090	3865	158	1929	Hoover	Herbert
##	1933-Roosevelt	743	2062	85	1933	Roosevelt	Franklin D.
##	1937-Roosevelt	725	1997	96	1937	Roosevelt	Franklin D.
##	1941-Roosevelt	526	1544	68	1941	Roosevelt	Franklin D.
##	1945-Roosevelt	275	647	26	1945	Roosevelt	Franklin D.
##	1949-Truman	781	2513	116	1949	Truman	Harry S.
##	1953-Eisenhower	900	2757	119	1953	Eisenhower	Dwight D.
##	1957-Eisenhower	621	1931	92	1957	Eisenhower	Dwight D.
##	1961-Kennedy	566	1566	52	1961	Kennedy	John F.
##	1965-Johnson	568	1723	93	1965	Johnson	Lyndon Baines
##	1969-Nixon	743	2437	103	1969	Nixon	Richard Milhous
##	1973-Nixon	544	2012	68	1973	Nixon	Richard Milhous
##	1977-Carter	527	1376	52	1977	Carter	Jimmy
##	1981-Reagan	902	2790	128	1981	Reagan	Ronald
##	1985-Reagan	925	2921	123	1985	Reagan	Ronald
##	1989-Bush	795	2681	141	1989	Bush	George
##	1993-Clinton	642	1833	81	1993	Clinton	Bill
##	1997-Clinton	773	2449	111	1997	Clinton	Bill
##	2001-Bush	621	1808	97	2001	Bush	George W.
##	2005-Bush	773	2319	100	2005	Bush	George W.
##	2009-Obama	938	2711	110	2009	Obama	Barack

```
##      2013-Obama   814   2317      88 2013      Obama      Barack
##      2017-Trump   582   1660      88 2017      Trump      Donald J.
##
## Source: Gerhard Peters and John T. Woolley. The American Presidency Project.
## Created: Tue Jun 13 14:51:47 2017
## Notes: http://www.presidency.ucsb.edu/inaugurals.php

toks <- tokens(full_inaug_corp, remove_punct=TRUE)

toks1 <- tokens_select(toks, stopwords('english'), selection='remove')

dfm1 <- dfm(toks1)
```

2. Lexical Diversity

There are times when we might be interested in the complexity of the words used in text. There are a ton of measures here, but they all roughly do the same thing - comparing the number of unique words (tokens) and total words (tokens).

```
lexdiv <- textstat_lexdiv(dfm1)
head(lexdiv, 5)
```

```
##      document      TTR      C      R      CTTR      U
## 1 1789-Washington 0.7806748 0.9617909 19.933978 14.095451 73.65389
## 2 1793-Washington 0.9354839 0.9838408 7.366007 5.208554 110.92051
## 3 1797-Adams      0.6542056 0.9391673 21.399624 15.131820 49.79857
## 4 1801-Jefferson 0.7293973 0.9529101 20.797418 14.705995 61.79862
## 5 1805-Jefferson 0.6726014 0.9426767 21.386186 15.122317 52.41763
##      S      Maas      lgV0      lgeV0
## 1 0.9623481 0.1165204 9.886289 22.76402
## 2 0.9720826 0.0949498 9.849038 22.67825
## 3 0.9433742 0.1417071 8.283634 19.07377
## 4 0.9548442 0.1272069 9.144363 21.05567
## 5 0.9463441 0.1381215 8.488001 19.54435
```


```
tail(lexdiv, 5)
```

```
##      document      TTR      C      R      CTTR      U      S
## 54 2001-Bush 0.6428571 0.9337026 18.00000 12.72792 43.65658 0.9354530
## 55 2005-Bush 0.6176753 0.9306568 19.92900 14.09193 43.51472 0.9349294
## 56 2009-Obama 0.6828645 0.9460250 23.38748 16.53745 56.86513 0.9505226
## 57 2013-Obama 0.6621622 0.9406254 21.31298 15.07055 50.78539 0.9445418
## 58 2017-Trump 0.6419580 0.9325604 17.16563 12.13793 42.32387 0.9334292
##      Maas      lgV0      lgeV0
## 54 0.1513475 7.547647 17.37910
## 55 0.1515940 7.674943 17.67221
## 56 0.1326102 8.959235 20.62940
## 57 0.1403236 8.355738 19.23980
## 58 0.1537118 7.373155 16.97732
```

We can also plot lexical diversity by selecting one of the measures. Below we do it with Type-Token Ratio, which is simple the number of unique words divided by the total number of words. (**Remember that if you are using a Mac, just skip the x11() function.**)

```
library(ggplot2)

x11()
ggplot(data=lexdiv, aes(x=document, y=TTR)) +
  geom_point() +
  coord_flip() +
  labs(x="President", y="Type-Token Ratio")
```



3. Keyness

Another somewhat simple analysis that is available is **keyness**. The basic idea is to compare words between a target and a reference group. Below we first compare pre-war (WWII), our reference group, and post-war, our target group, speeches.

```
period <- ifelse(docvars(data_corpus_inaugural, "Year") < 1945, "pre-war", "post-war")
```

```
dfm2 <- dfm(toks1, groups = period)
head(dfm2)
```

```
## Document-feature matrix of: 2 documents, 9,205 features (35.4% sparse).
```

```
head(result <- textstat_keyness(dfm2), 10)
```

```
##      feature      chi2 p n_target n_reference
## 1      us 186.34621 0      262      216
```

```
## 2    america 176.58663 0      130      54
## 3      world 175.24439 0      188     123
## 4  americans 150.53636 0       67       7
## 5       new 141.37882 0      150      97
## 6   freedom 137.30596 0      121      64
## 7    today 128.77885 0       76      21
## 8      let 111.45125 0      100      54
## 9   together 104.88823 0       64      19
## 10 america's 93.97685 0       35       0
```

```
tail(result <- textstat_keyness(dfm2), 10)
```

```
##           feature      chi2      p n_target n_reference
## 9196      state -25.85907 3.672751e-07      6      103
## 9197    powers -27.49893 1.571812e-07      4       97
## 9198  executive -28.42150 9.757587e-08      3       94
## 9199      may -30.12530 4.050120e-08     47      291
## 9200    policy -33.50265 7.116694e-09      1       96
## 9201  government -43.19681 4.950129e-11     84      480
## 9202      upon -52.21073 4.984901e-13     39      332
## 9203    public -56.19816 6.550316e-14     11      213
## 9204    states -59.43006 1.265654e-14     28      305
## 9205 constitution -61.37027 4.773959e-15      6      200
```

The two right-hand columns are the number of times the words appear between the two groups. The Chi-Square value is signed positively if the words appear greater than expected and negative if vice-versa. The *p*-values seem essentially worthless since they seem like they are always ‘statistically significant’.

There are several different options for the statistics reported with `textstat_keyness()`. Let’s look at the commonly utilised odds-ratio by including `measure="exact"`. The one problem for very sparse dfms is that the comparisons may be against 0, which makes for infinitely greater odds or 0 odds.

```
head(result1 <- textstat_keyness(dfm2, measure="exact"), 10)
```

```
##           feature or      p n_target n_reference
## 1    america's Inf 1.478186e-20      35       0
## 2    faintness Inf 2.714737e-01       1       0
## 3 schoolmaster Inf 2.714737e-01       1       0
## 4      dr Inf 7.369489e-02       2       0
## 5    peabody Inf 2.714737e-01       1       0
## 6 untroubled Inf 2.714737e-01       1       0
## 7    smoothly Inf 2.714737e-01       1       0
## 8    downward Inf 2.714737e-01       1       0
## 9      trend Inf 2.000456e-02       3       0
## 10    peaks Inf 7.369489e-02       2       0
```


```
tail(result1 <- textstat_keyness(dfm2, measure="exact"), 10)
```

```
##           feature or p n_target n_reference
## 9196  unnoticed 0 1      0      1
## 9197      48 0 1      0      1
## 9198   counties 0 1      0      1
## 9199    towns 0 1      0      1
## 9200   prophecy 0 1      0      1
## 9201    1941 0 1      0      1
## 9202  smothered 0 1      0      1
## 9203    strove 0 1      0      1
```

```
## 9204 valiantly 0 1 0 1
## 9205 triumphantly 0 1 0 1
```

Finally, we can do a plot using `textplot_keyness()` for keyness of the words with the largest chi-square values for the target and reference groups.

```
x11()
textplot_keyness(result)
```



4. Document Similarity

We may be interested in the similarity of documents in documents. To get basic correlations between documents, we can use the `textstat_simil()` function. Since the correlation matrix will be huge, let's subset the corpus to only include speeches since 1992.

```
presDfm <- dfm(corpus_subset(data_corpus_inaugural, Year > 1992),
               remove = stopwords("english"), stem = TRUE, remove_punct = TRUE)

simil1 <- textstat_simil(presDfm, margin="documents")
simil1
```

```
##           1993-Clinton 1997-Clinton 2001-Bush 2005-Bush 2009-Obama
## 1997-Clinton    0.6271129
## 2001-Bush       0.4734439    0.5327868
## 2005-Bush       0.4338601    0.4289271 0.5344415
## 2009-Obama      0.5707623    0.6026942 0.5241995 0.4330978
```

```
## 2013-Obama      0.5725041    0.5916516 0.5523905 0.5137096 0.6103693
## 2017-Trump      0.4967910    0.4989669 0.4672634 0.4739241 0.4377484
##
## 2013-Obama
## 1997-Clinton
## 2001-Bush
## 2005-Bush
## 2009-Obama
## 2013-Obama
## 2017-Trump      0.4414144
```

This gives the correlations between each of the speeches since 1992. As with most of the `Quanteda` functions, there are a bunch of different options for statistics to examine.


5. Locations of Keywords in Documents

If we are interested in where certain words (i.e., keywords) are located in documents, we can use the `textplot_xray()` function. Because we are considering the locations of certain words we need to work with the corpus and not the `dfm`, and we need to also need to use the `kwic()` function. Let's first look at the words 'american' and 'people'.

```
data_corpus_inaugural_subset <- corpus_subset(data_corpus_inaugural, Year > 1992)

x11()
textplot_xray(
  kwic(data_corpus_inaugural_subset, "american"),
  kwic(data_corpus_inaugural_subset, "people")
)
```

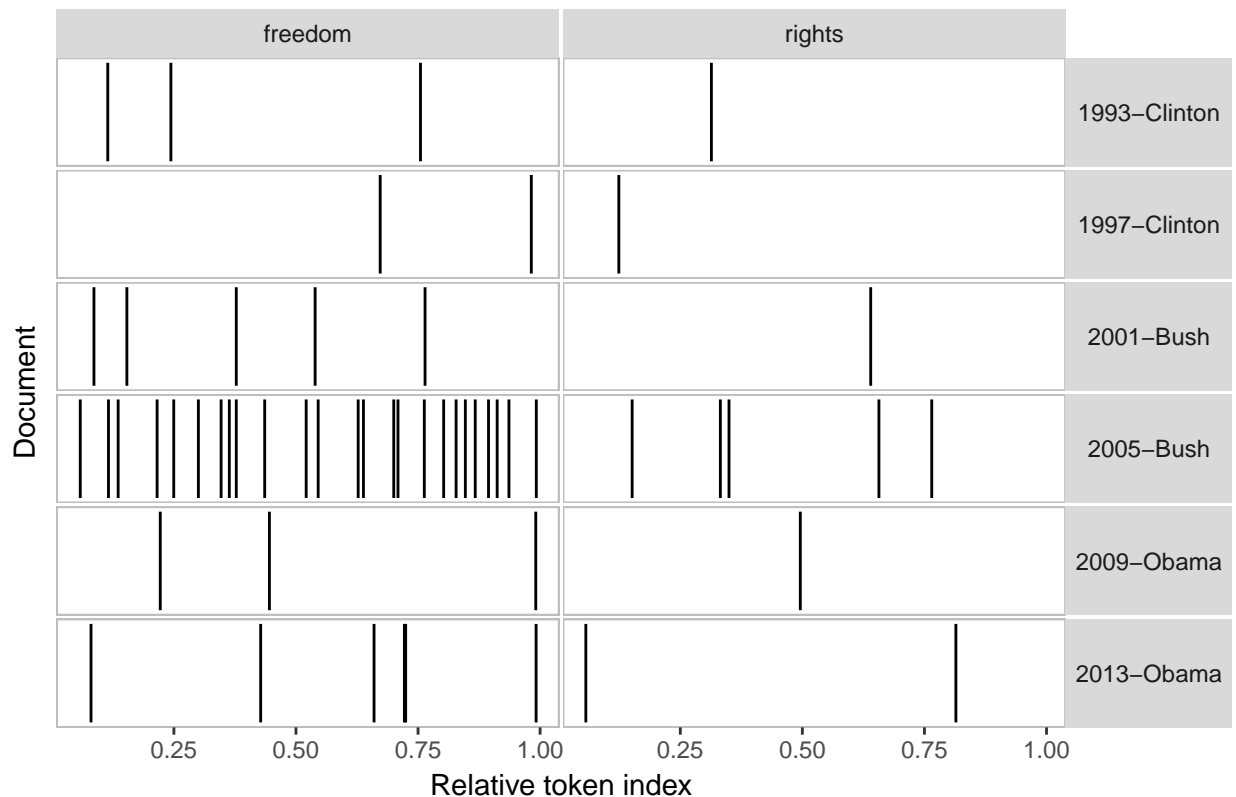
Lexical dispersion plot



Now let's try 'freedom' and 'rights'.

```
x11()
textplot_xray(
  kwic(data_corpus_inaugural_subset, "freedom"),
  kwic(data_corpus_inaugural_subset, "rights")
)
```


Lexical dispersion plot



(Note: Trump does not appear because it seems he did not talk about ‘freedom’ or ‘rights’.)

6. Sentiment Analysis

Looking at the sentiment of text (how positive or negative) is a common type of analysis done in social science. As with all the different types of text analysis, whether or not to look at sentiment is dependent on your research question.

We are going to use the **Quanteda** built-in Lexicoder Sentiment Dictionary (LSD) developed by Young and Soroka.

```
lengths(data_dictionary_LSD2015)
```

```
##      negative      positive neg_positive neg_negative
##      2858        1709        1721        2860
```

This shows the number of words that are classified as negative and positive in the dictionary. The **neg_positive** column are positive words that are preceded by a negation (e.g., ‘not good’) and thus are negative. The **neg_negative** column are negative words that are preceded by a negation (e.g., ‘not bad’) and thus are positive. Words that are not expressly positive or negative (‘neutral’ words) are excluded here. To create a sentiment dfm, we use the **dfm_lookup()** function.

```
lsd_dfm <- dfm_lookup(dfm1, data_dictionary_LSD2015)
head(lsd_dfm)
```

```
## Document-feature matrix of: 6 documents, 4 features (50% sparse).
## 6 x 4 sparse Matrix of class "dfm"
##           features
```

```
## docs          negative positive neg_positive neg_negative
## 1789-Washington 39      121          0          0
## 1793-Washington 3       10          0          0
## 1797-Adams      60     238          0          0
## 1801-Jefferson  66     177          0          0
## 1805-Jefferson  85     164          0          0
## 1809-Madison    55     138          0          0
```

```
tail(lsd_dfm)
```

```
## Document-feature matrix of: 6 documents, 4 features (50% sparse).
```

```
## 6 x 4 sparse Matrix of class "dfm"
```

```
##          features
## docs          negative positive neg_positive neg_negative
## 1997-Clinton    69      147          0          0
## 2001-Bush        62      151          0          0
## 2005-Bush        86      212          0          0
## 2009-Obama      111      192          0          0
## 2013-Obama       76      169          0          0
## 2017-Trump       43      123          0          0
```

Even though Obama appears to have more words that are classified, it is surprising how many are classified as negative.

Often we prefer to plot out the sentiment values. Because `ggplot()` does not work with `dfms`, we will transform the `dfm` to a standard data frame using the `as.data.frame()` function. We first plot the negative words used.


```
lsd <- as.data.frame(lsd_dfm)
```

```
lsd
```

```
##          document negative positive neg_positive neg_negative
## 1 1789-Washington 39      121          0          0
## 2 1793-Washington 3       10          0          0
## 3 1797-Adams      60     238          0          0
## 4 1801-Jefferson  66     177          0          0
## 5 1805-Jefferson  85     164          0          0
## 6 1809-Madison    55     138          0          0
## 7 1813-Madison    86     106          0          0
## 8 1817-Monroe     110     338          0          0
## 9 1821-Monroe     144     344          0          0
## 10 1825-Adams      95     242          0          0
## 11 1829-Jackson    28     110          0          0
## 12 1833-Jackson    32      90          0          0
## 13 1837-VanBuren   180     351          0          0
## 14 1841-Harrison   280     569          0          0
## 15 1845-Polk       166     388          0          0
## 16 1849-Taylor      19     127          0          0
## 17 1853-Pierce     114     331          0          0
## 18 1857-Buchanan   111     247          0          0
## 19 1861-Lincoln    157     209          0          0
## 20 1865-Lincoln     42      37          0          0
## 21 1869-Grant      31      97          0          0
## 22 1873-Grant      33      95          0          0
## 23 1877-Hayes      62     218          0          0
## 24 1881-Garfield   119     233          0          0
## 25 1885-Cleveland  52     176          0          0
```

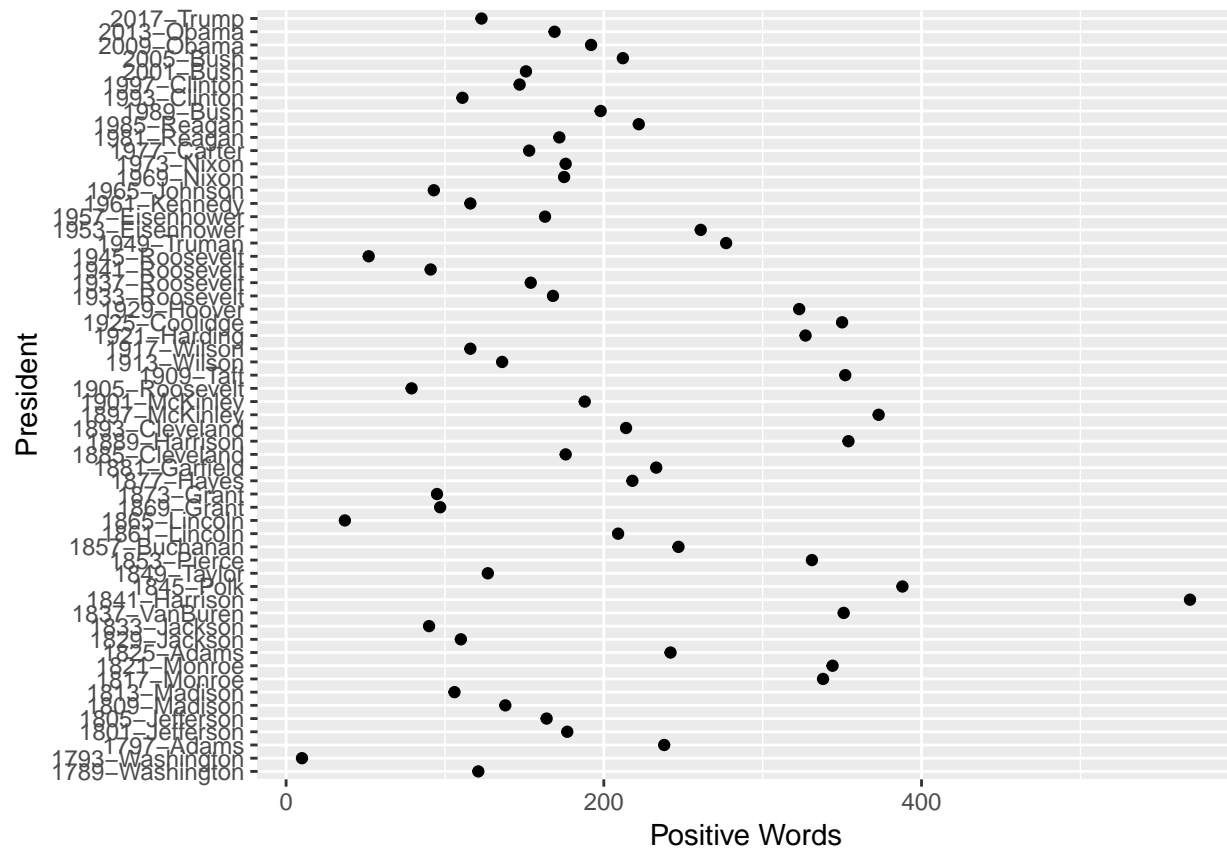
## 26	1889-Harrison	156	354	0	0
## 27	1893-Cleveland	116	214	0	0
## 28	1897-McKinley	111	373	0	0
## 29	1901-McKinley	73	188	0	0
## 30	1905-Roosevelt	37	79	0	0
## 31	1909-Taft	154	352	0	0
## 32	1913-Wilson	71	136	0	0
## 33	1917-Wilson	44	116	0	0
## 34	1921-Harding	163	327	0	0
## 35	1925-Coolidge	125	350	0	0
## 36	1929-Hoover	139	323	0	0
## 37	1933-Roosevelt	88	168	0	0
## 38	1937-Roosevelt	81	154	0	0
## 39	1941-Roosevelt	36	91	0	0
## 40	1945-Roosevelt	17	52	0	0
## 41	1949-Truman	76	277	0	0
## 42	1953-Eisenhower	92	261	0	0
## 43	1957-Eisenhower	68	163	0	0
## 44	1961-Kennedy	80	116	0	0
## 45	1965-Johnson	60	93	0	0
## 46	1969-Nixon	74	175	0	0
## 47	1973-Nixon	49	176	0	0
## 48	1977-Carter	41	153	0	0
## 49	1981-Reagan	73	172	0	0
## 50	1985-Reagan	84	222	0	0
## 51	1989-Bush	51	198	0	0
## 52	1993-Clinton	61	111	0	0
## 53	1997-Clinton	69	147	0	0
## 54	2001-Bush	62	151	0	0
## 55	2005-Bush	86	212	0	0
## 56	2009-Obama	111	192	0	0
## 57	2013-Obama	76	169	0	0
## 58	2017-Trump	43	123	0	0

```
x11()
ggplot(data=lsd, aes(x=document, y=negative)) +
  geom_point() +
  coord_flip() +
  labs(x="President", y="Negative Words")
```



Now, we plot the positive words used.


```
x11()
ggplot(data=lsd, aes(x=document, y=positive)) +
  geom_point() +
  coord_flip() +
  labs(x="President", y="Positive Words")
```



Since the previous plots are a bit misleading because it is just the raw counts, let's do a plot of how positive or negative the speech was out of the total number of positive and negative words used.

```
lsd$overall <- (lsd$positive-lsd$negative)/(lsd$positive+lsd$negative)
```

```
x11()
ggplot(data=lsd, aes(x=document, y=overall)) +
  geom_point() +
  coord_flip() +
  labs(x="President",y="Overall Sentiment")
```



Interestingly, Lincoln’s 1865 speech is the only one that had more negative than positive words; makes some sense since it was the last year of the US Civil War.

7. Topic Models

Topic models in general are increasingly used by social scientists in situations where a researcher wants to uncover themes or topics within the text. There are a variety of approaches researchers have used (Grimmer 2010; Grimmer 2013), but the Latent Dirichlet Allocation (LDA) approach is one of the simplest and most widely employed (Blei et al. 2003; Blei 2012). The basic concept of the LDA process is to assume unobserved topics exist amongst a set of documents; the documents are observable, but the topics are latent. The assumption is that the topics arise probabilistically from the words in the documents following a Dirichlet distribution. Different topics have different words that arise with higher probabilities than other words. For example, a ‘free speech’ topic will have words associated with free speech at higher probabilities than non-free speech words. This unsupervised, algorithmic approach is an efficient way of ascertaining meaning and structure out of a large set of text.

Topic models are not part of **Quanteda**’s internal package, but they are still simple to do using **Quanteda**. To do so, we use the `convert()` function to create a dtm from the dfm. Then we use the `LDA()` function from the **topicmodels** package and ask for 10 topics ($k = 10$). The number of topics is subjective and is more-or-less an art dependent on the size of the corpus; though 10 topics is usually the standard starting point. Running LDA topic models is computationally intensive where the larger the size of the corpus/dfm/dtm, the longer it will take.

```
library(topicmodels)
```

```
## Warning: package 'topicmodels' was built under R version 3.4.4
```

```
dtm <- convert(dfm1, to="topicmodels")
lda <- LDA(dtm, k=10)
```

```
terms(lda,10)
```

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5
## [1,]	"people"	"government"	"us"	"upon"	"us"
## [2,]	"shall"	"great"	"world"	"government"	"new"
## [3,]	"can"	"states"	"freedom"	"people"	"world"
## [4,]	"upon"	"country"	"can"	"may"	"people"
## [5,]	"may"	"people"	"must"	"every"	"america"
## [6,]	"states"	"public"	"peace"	"can"	"must"
## [7,]	"constitution"	"every"	"people"	"us"	"nation"
## [8,]	"government"	"united"	"nation"	"states"	"can"
## [9,]	"laws"	"union"	"nations"	"public"	"time"
## [10,]	"now"	"may"	"new"	"country"	"let"

	Topic 6	Topic 7	Topic 8	Topic 9	Topic 10
## [1,]	"government"	"people"	"people"	"government"	"can"
## [2,]	"states"	"government"	"government"	"upon"	"great"
## [3,]	"union"	"states"	"must"	"people"	"upon"
## [4,]	"may"	"war"	"can"	"can"	"government"
## [5,]	"country"	"nation"	"world"	"may"	"people"
## [6,]	"people"	"upon"	"shall"	"power"	"country"
## [7,]	"shall"	"can"	"us"	"must"	"every"
## [8,]	"one"	"shall"	"may"	"states"	"must"
## [9,]	"constitution"	"may"	"public"	"country"	"nation"
## [10,]	"can"	"united"	"national"	"great"	"world"

As should be apparent from the output, the substantive meaning of each topic is often not readily clear. This will vary depending on the substance of our corpus. For example, if our corpus has fairly similar documents/texts, then differences will be hard to discern. If our corpus has very different documents/texts, then the topic differences will be clearer. Here, in particular, the meaning of each topic is not clear; particularly since the same words can appear in different topics. Hence, we might want to re-run the model and look at only 5 topics.

```
lda2 <- LDA(dtm, k=5)
```

```
terms(lda2,10)
```

	Topic 1	Topic 2	Topic 3	Topic 4	Topic 5
## [1,]	"can"	"government"	"government"	"people"	"us"
## [2,]	"us"	"people"	"can"	"government"	"must"
## [3,]	"world"	"states"	"people"	"upon"	"world"
## [4,]	"peace"	"may"	"upon"	"states"	"new"
## [5,]	"must"	"great"	"must"	"public"	"america"
## [6,]	"people"	"country"	"states"	"shall"	"people"
## [7,]	"shall"	"power"	"may"	"country"	"nation"
## [8,]	"great"	"upon"	"world"	"may"	"freedom"
## [9,]	"new"	"union"	"peace"	"can"	"can"
## [10,]	"government"	"constitution"	"now"	"great"	"time"

Even with only 5 topics, the differences in topics is not evident.

We can also look at the most likely topic for each document using the `topics()` function.

```
head(topics(lda),20)
```

```
## 1789-Washington 1793-Washington      1797-Adams 1801-Jefferson
##              10              1              7              6
## 1805-Jefferson 1809-Madison      1813-Madison 1817-Monroe
##              4              2              7              2
## 1821-Monroe   1825-Adams      1829-Jackson 1833-Jackson
##              2              2              2              4
## 1837-VanBuren 1841-Harrison      1845-Polk 1849-Taylor
##              4              9              6              1
## 1853-Pierce 1857-Buchanan      1861-Lincoln 1865-Lincoln
##              4              6              1              7
```

```
tail(topics(lda),20)
```

```
## 1941-Roosevelt 1945-Roosevelt      1949-Truman 1953-Eisenhower
##              3              3              3              3
## 1957-Eisenhower 1961-Kennedy      1965-Johnson 1969-Nixon
##              3              3              5              5
## 1973-Nixon      1977-Carter      1981-Reagan 1985-Reagan
##              3              3              3              5
## 1989-Bush      1993-Clinton      1997-Clinton 2001-Bush
##              5              5              5              3
## 2005-Bush      2009-Obama      2013-Obama 2017-Trump
##              3              5              3              5
```

8. Ideological Placements (Scalings)

One particularly large area of political science deals with ideological placements of parties, candidates, politicians, interest groups, etc. Historically, this meant using votes (e.g., DW-Nominate), ‘experts’ (e.g., Chapel Hill Expert Survey), or the subject themselves to create ideological placements. Increasingly, researchers are using text (e.g., party platforms, speeches, content of bills, etc.) to make ideological placements; moving from what actors do to what actors do and say.

Today, we look at two varieties of placements using **Wordscores** and **Wordfish**. Wordscores has been around for over decade, while Wordfish is much newer. Wordscores requires reference texts that says what is ‘liberal’ and ‘conservative’; or whatever version one is interested in. Then we compare all remaining texts to the reference texts. Therefore, we can consider Wordscores a **supervised method**. Therefore, if your research does not have legitimate reference texts then you cannot use Wordscores. For example, we cannot use either approach for creating ideological placements of newspapers. Why? Because we do not have clear reference texts as to what is conservative and what is liberal.

Wordfish meanwhile can be considered an **unsupervised method** method because it does not rely on reference texts. Instead of relying on the reference texts, it estimates documents positions solely based on the frequency of the words. So, can we use Wordfish for ideological placements/scalings of any type of document? According to Slapin (co-author of Wordfish), no. Even though Wordfish does not rely on reference texts, Slapin still is very cautious about using Wordfish when working with documents that do not clearly have ideological differences. For example, Slapin says we cannot use Wordfish to place newspapers... unless the newspapers are owned/controlled by political parties.

8.1. Wordscores

Due to the requirement to need a reference, we will make use of Quanteda's built-in corpus of 2010 Irish budget speeches (data_corpus_irishbudget2010).

```
ie_dfm <- data_corpus_irishbudget2010 %>%  
  dfm(remove = stopwords("english"), remove_punct = TRUE)  
ie_dfm
```

```
## Document-feature matrix of: 14 documents, 5,008 features (82.7% sparse).
```

```
# Set reference scores  
refscores <- c(rep(NA, 4), 1, -1, rep(NA, 8))  
  
# Predict Wordscores model  
ws <- textmodel_wordscores(ie_dfm, refscores, smooth = 1)  
ws
```

```
##  
## Call:  
## textmodel_wordscores.dfm(x = ie_dfm, y = refscores, smooth = 1)  
##  
## Scale: linear; 2 reference scores; 5008 scored features.  
summary(ws)
```

```
##  
## Call:  
## textmodel_wordscores.dfm(x = ie_dfm, y = refscores, smooth = 1)  
##  
## Reference Document Statistics:  
##
```

	score	total	min	max	mean	median
## 2010_BUDGET_01_Brian_Lenihan_FF	NA	9286	1	39	1.854	1
## 2010_BUDGET_02_Richard_Bruton_FG	NA	6865	1	46	1.371	1
## 2010_BUDGET_03_Joan_Burton_LAB	NA	7976	1	42	1.593	1
## 2010_BUDGET_04_Arthur_Morgan_SF	NA	8219	1	54	1.641	1
## 2010_BUDGET_05_Brian_Cowen_FF	1	8205	1	34	1.638	1
## 2010_BUDGET_06_Enda_Kenny_FG	-1	6879	1	28	1.374	1
## 2010_BUDGET_07_Kieran_ODonnell_FG	NA	6011	1	26	1.200	1
## 2010_BUDGET_08_Eamon_Gilmore_LAB	NA	6934	1	29	1.385	1
## 2010_BUDGET_09_Michael_Higgins_LAB	NA	5526	1	9	1.103	1
## 2010_BUDGET_10_Ruairi_Quinn_LAB	NA	5522	1	13	1.103	1
## 2010_BUDGET_11_John_Gormley_Green	NA	5485	1	11	1.095	1
## 2010_BUDGET_12_Eamon_Ryan_Green	NA	5692	1	15	1.137	1
## 2010_BUDGET_13_Ciaran_Cuffe_Green	NA	5582	1	19	1.115	1
## 2010_BUDGET_14_Caoimhghin_OCaolain_SF	NA	6911	1	29	1.380	1

```
##  
## Wordscores:  
## (showing first 30 elements)  
##
```

	presented	supplementary	budget	house
##	-0.40925	-0.40925	-0.13089	0.61479
##	last	april	said	work
##	0.07993	-0.71280	-0.71280	-0.02137
##	way	period	severe	economic
##	0.01220	0.35400	0.25283	0.34126
##	distress	today	can	report

We see that two of the texts are set as the reference points - one at 1 and the other at -1. Next, we plot the word position and highlight a few in red.

[illegible]

```
pred <- predict(ws)
pred
```

18


```
##      2010_BUDGET_03_Joan_Burton_LAB
##      -0.03005066
##      2010_BUDGET_04_Arthur_Morgan_SF
##      -0.01329439
##      2010_BUDGET_05_Brian_Cowen_FF
##      0.10685927
##      2010_BUDGET_06_Enda_Kenny_FG
##      -0.10685927
##      2010_BUDGET_07_Kieran_ODonnell_FG
##      -0.03316757
##      2010_BUDGET_08_Eamon_Gilmore_LAB
##      -0.02613156
##      2010_BUDGET_09_Michael_Higgins_LAB
##      -0.03679662
##      2010_BUDGET_10_Ruairi_Quinn_LAB
##      -0.03792328
##      2010_BUDGET_11_John_Gormley_Green
##      -0.02334390
##      2010_BUDGET_12_Eamon_Ryan_Green
##      -0.02379303
##      2010_BUDGET_13_Ciaran_Cuffe_Green
##      -0.02469264
##      2010_BUDGET_14_Caoimhghin_OCaolain_SF
##      -0.02512630
```

This is now the placements of the budget speeches.

We can also plot the placements of the budget speeches using the `textplot_scale1d()`. First we set the document labels for the y-axis and then we plot the predicted placements.

```
# Set document labels for y-axis
doclab <- apply(docvars(data_corpus_irishbudget2010, c("name", "party")),
               1, paste, collapse = " ")


# Plot estimated document positions and group by "party" variable
x11()
textplot_scale1d(pred, margin = "documents",
                 doclabels = doclab,
                 groups = docvars(data_corpus_irishbudget2010, "party"))
```



We can re-scale the x-axis to make a clearer comparison with Wordfish below.

```
pred_lbg <- predict(ws, rescaling = "lbg")

x11()
textplot_scale1d(pred_lbg, margin = "documents",
  doclabels = doclab,
  groups = docvars(data_corpus_irishbudget2010, "party"))
```



8.2. Wordfish

Now let's replicate the placements using the `textmodel_wordfish()` function with `ie_dfm`. There are a few different ways to identify the Wordfish model. Below we will follow the [Quanteda](#) tutorial by fixing two texts and specifying that text 6 is to the left of text 5 (`dir=c(6,5)`); though note that doing this does not set the scaling domain like Wordscores does.

```
wf <- textmodel_wordfish(ie_dfm, dir = c(6,5))
```

```
wf
```

```
##
```

```
## Call:
```

```
## textmodel_wordfish.dfm(x = ie_dfm, dir = c(6, 5))
```

```
##
```

```
## Dispersion: poisson; direction: 6 < 5; 14 documents; 5008 features.
```

```
summary(wf)
```

```
##
```

```
## Call:
```

```
## textmodel_wordfish.dfm(x = ie_dfm, dir = c(6, 5))
```

```
##
```

```
## Estimated Document Positions:
```

```
##
```

	theta	se
## 2010_BUDGET_01_Brian_Lenihan_FF	1.726667	0.02238
## 2010_BUDGET_02_Richard_Bruton_FG	-0.497193	0.03087

```

## 2010_BUDGET_03_Joan_Burton_LAB      -0.991714 0.01804
## 2010_BUDGET_04_Arthur_Morgan_SF      0.085603 0.02889
## 2010_BUDGET_05_Brian_Cowen_FF        1.903242 0.02479
## 2010_BUDGET_06_Enda_Kenny_FG        -0.776897 0.02635
## 2010_BUDGET_07_Kieran_ODonnell_FG    -0.283465 0.04612
## 2010_BUDGET_08_Eamon_Gilmore_LAB     -0.378940 0.03201
## 2010_BUDGET_09_Michael_Higgins_LAB   -1.210983 0.03562
## 2010_BUDGET_10_Ruairi_Quinn_LAB      -1.231395 0.03503
## 2010_BUDGET_11_John_Gormley_Green     1.039543 0.07677
## 2010_BUDGET_12_Eamon_Ryan_Green       0.114674 0.06288
## 2010_BUDGET_13_Ciaran_Cuffe_Green     0.504108 0.07199
## 2010_BUDGET_14_Caoimhghin_OCaolain_SF -0.003249 0.03671
##
## Estimated Feature Scores:
##      presented supplementary budget house last april said work
## beta   0.3181          1.092 0.04255 0.03628 0.2550 -0.2138 -0.9591 0.5193
## psi   -1.8281          -1.184 2.68012 1.01149 0.9405 -0.5944 -0.4582 1.0680
##      way period severe economic distress today can report
## beta 0.2949 0.5526 1.316 0.4712 1.704 0.1071 0.3369 0.6585
## psi 1.3705 -0.2549 -2.151 1.5029 -4.316 0.8039 1.5100 -0.3204
##      notwithstanding difficulties past eight months now road
## beta          1.704          1.231 0.5029 1.704 0.7138 0.3265 0.1144
## psi          -4.316          -1.455 0.8701 -4.316 0.2002 1.5366 -0.1131
##      recovery enormous benefit main political parties share
## beta 0.3986 -0.05312 -0.04021 0.9277 -0.3670 0.5242 -0.09684
## psi 0.8296 -1.08890 1.33930 -0.9169 -0.4442 -0.1218 -0.17513

```

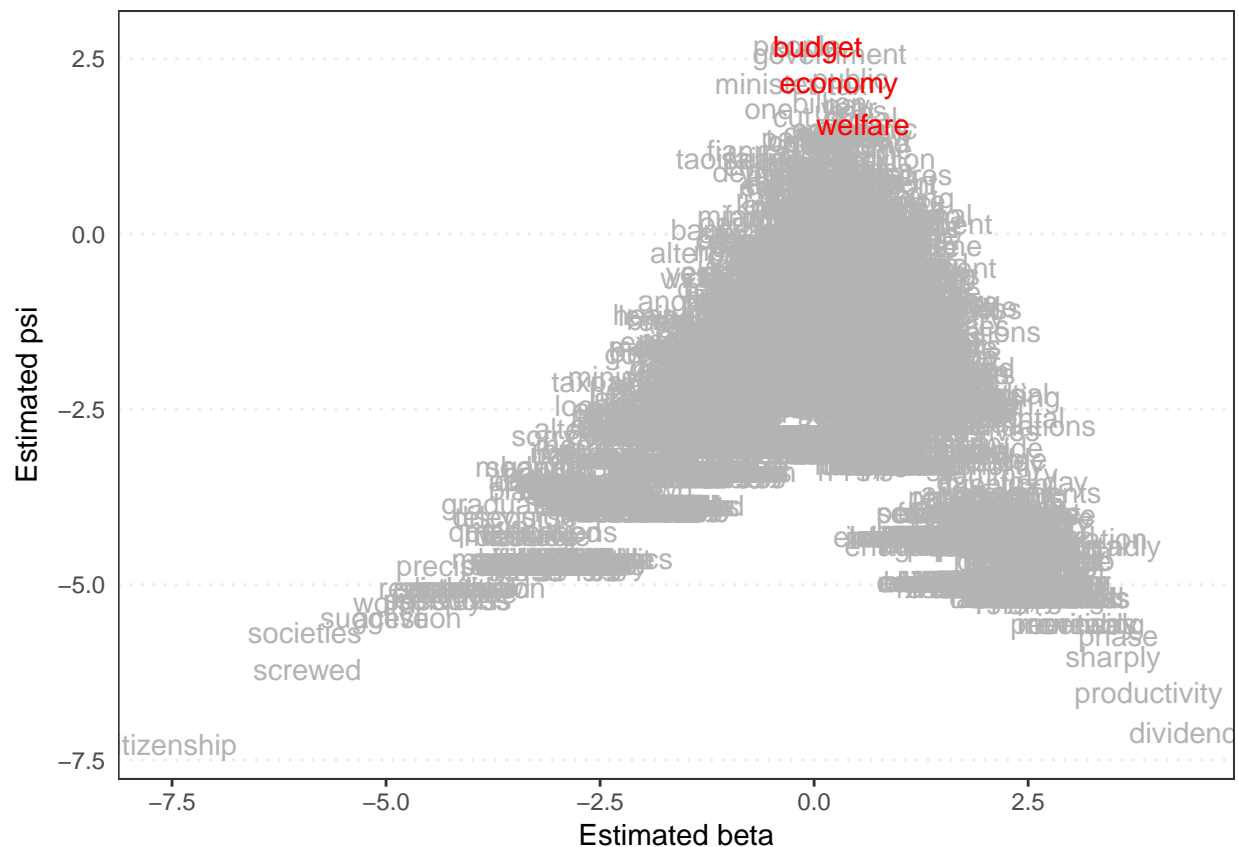
There are three things to note about the output. First, **theta** provides the estimated document positions; again, positive means conservative and negative means liberal. Second, for the estimates feature scores, **beta** is the word weight (informativeness) and **psi** is the word frequency; these are normalised to 0, thus we can have negative values.

Next, we plot β and ψ and highlight a few in red.


```

x11()
textplot_scale1d(wf, margin="features",
                 highlighted = c("economy", "welfare", "budget"),
                 highlighted_color = "red")

```



Unlike Wordscores, we do not need to run a prediction function to estimate the document positions, because Wordfish has already done that. Now let's plot the documents.



We see that compared to Wordscores, the Wordfish placements are similar and also automatically provide confidence intervals.

9. On Your Own

Try working through the above with your own data or the voter fraud data. For the voter fraud data, without additional manipulation, you can only look at lexical diversity, sentiment, and topic models.